


What aggregation recovers: performance against data heterogeneity

Macro-F1 on the shared held-out test set ($n = 161$), identical partitions for both training regimes



Label-distribution heterogeneity across the 5 clients (mean Jensen-Shannon divergence) →

—●— Federated (FedAvg) —●— Isolated (local-only) — — Centralized ceiling

seed 42 · split cb67241e1200 · FedAvg: 10 rounds × 3 local epochs × 5 clients · MLP 20264-512-128-5